

# **Unit 9: Transactions**

# Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
- E.g. transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- Two main issues to deal with:
  - **Failures of various kinds**, such as hardware failures and system crashes while executing transactions.
  - **Concurrent execution** of multiple transactions without inconsistency.

# ACID Properties

A **transaction** is a unit of program execution that accesses and possibly updates various data items.

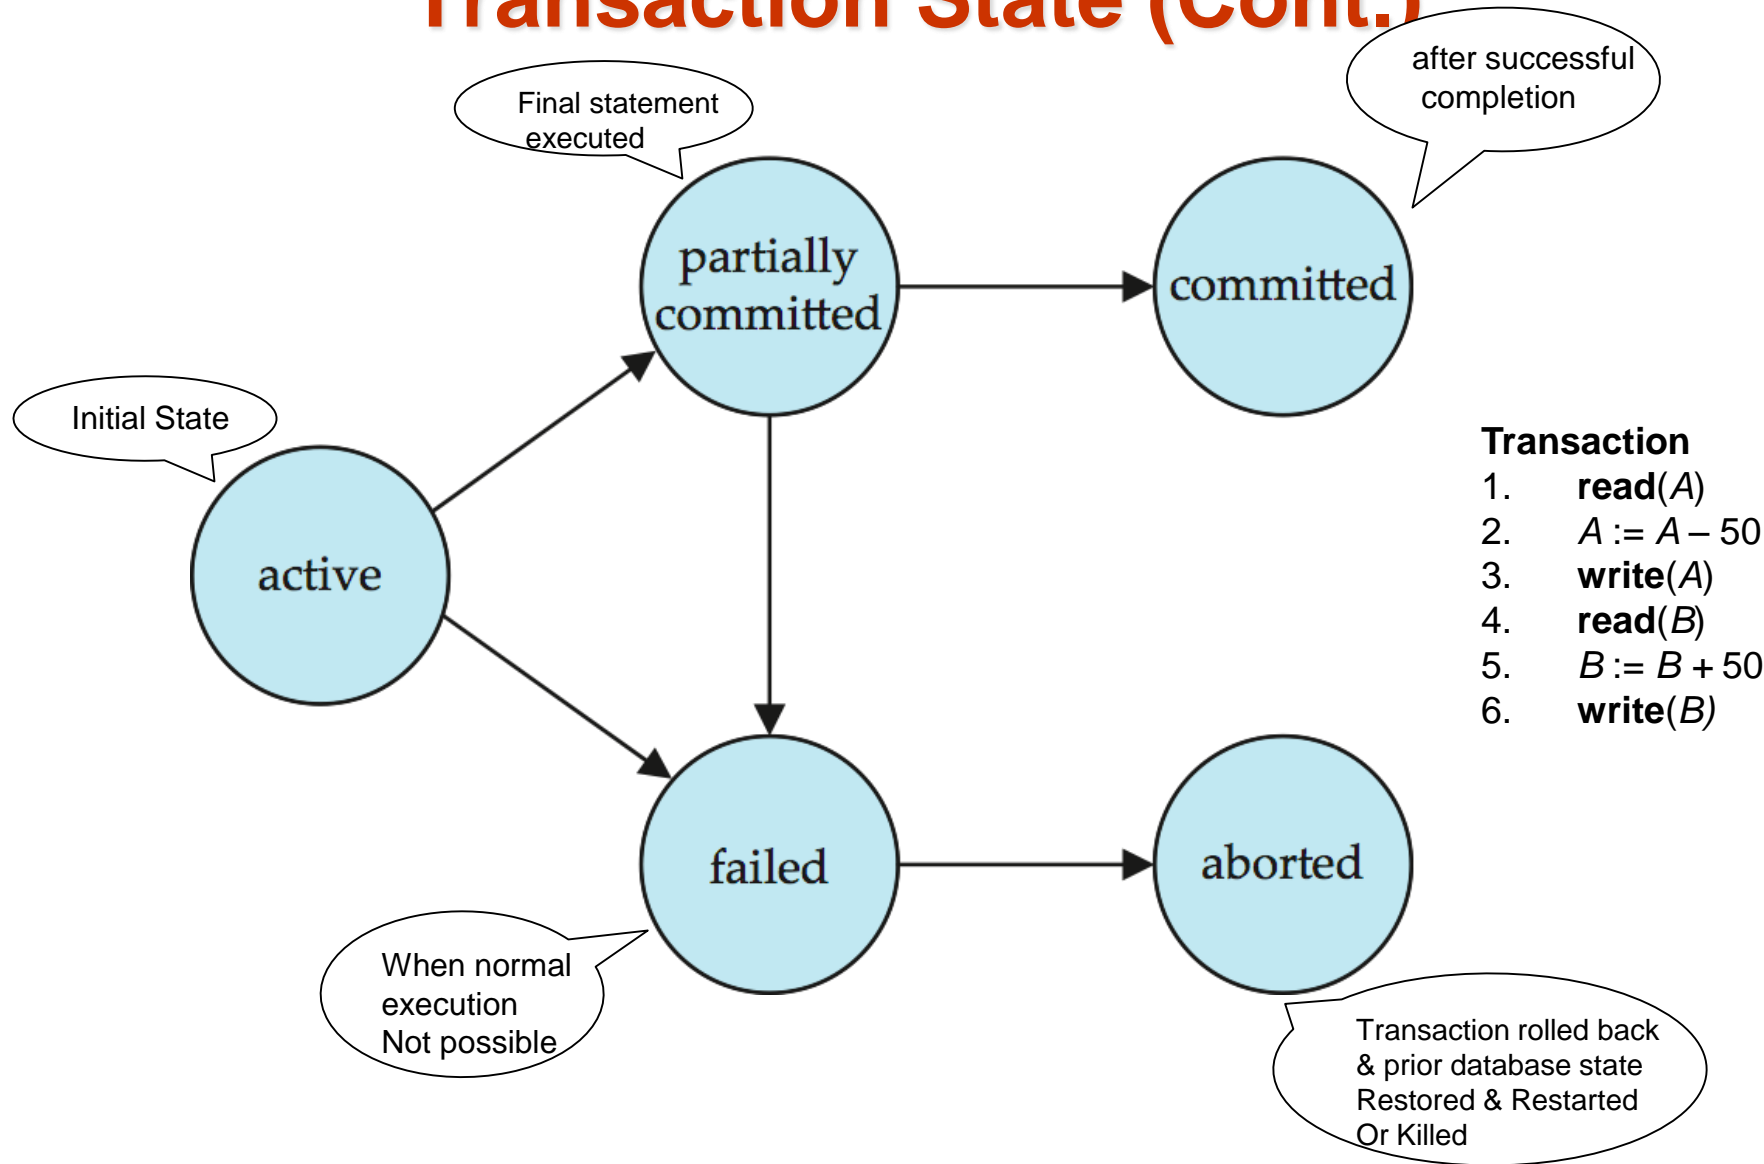
**To preserve the integrity of data** the database system must ensure:

- ❑ **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
- ❑ **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- ❑ **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. **Intermediate transaction results must be hidden from other concurrently executed transactions.**
- ❑ **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

# Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** – after the discovery that normal execution can no longer proceed.
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
  - restart the transaction
    - ▶ can be done only if no internal logical error
  - kill the transaction
- **Committed** – after successful completion.

# Transaction State (Cont.)



# Schedules

- **Schedule** – A **sequences of instructions** that specify the **chronological order** in which instructions of concurrent transactions are executed.
  - a schedule for a set of transactions must consist of **all instructions** of those transactions.
  - must preserve the **order** in which the instructions appear in each individual transaction.
- A transaction **that successfully completes** its execution will have a **commit instructions as the last statement**
  - by default transaction assumed to execute commit instruction as its last step
- A transaction **that fails** to successfully complete its execution will have an **abort instruction as the last statement**

# Schedule 1 (Serial)

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- Suppose the current values of accounts  $A$  and  $B$  are \$1000 and \$2000.
- A **serial schedule** in which  $T_1$  is followed by  $T_2$  :

The final values of accounts  $A$  and  $B$ , after the execution in Schedule 1 takes place, are \$855 and \$2145, respectively.

Thus, the total amount of money in accounts  $A$  and  $B$ —that is, the **sum  $A + B$** —is **preserved** after the execution of both transactions.

At the end of **T1**  $A=950$  and  $B=2050$

At the end of **T2**  $A=855$  and  $B=2145$

Total Amount in the system is  $2145+855=3000$

Database **consistency maintained.**

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

## Schedule 2 (Serial)

- A **serial schedule** where  $T_2$  is followed by  $T_1$

Assume current value of  $A=1000$  and  $B=2000$

Similarly, if the transactions are executed one at a time in the order  $T_2$  followed by  $T_1$ , then the corresponding execution sequence is in schedule 2

Again, as expected, the **sum  $A + B$  is preserved**, and the final values of accounts  $A$  and  $B$  are **\$850** and **\$2150**, respectively.

At the end of  **$T_2$**   **$A=900$**  and  **$B=2100$**

At the end of  **$T_1$**   **$A=850$**  and  **$B=2150$**

Total Amount in the system is  $2150+850=3000$

Database **consistency maintained**.

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	



## Schedule 3 (Concurrent)

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is **not a serial schedule**, but it is *equivalent* to Schedule 1.

**accounts A and B are \$1000 and \$2000.**

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Finally  $A=855$  and  $B=2145$   
;  $A+B=855+2145=3000$   
—**consistency maintained**

In Schedules 1, 2 and 3, the **sum A + B is preserved**.

# Schedule 4 (Concurrent)

- The following concurrent schedule does not preserve the value of  $(A + B)$ . (Assume initially  $A=1000/-$  and  $B=2000/-$ )

	$T_1$	$T_2$	
A=1000	read (A)		A =1000
1000-50=950	$A := A - 50$		Temp=100
		read (A)	1000-100=900
		$temp := A * 0.1$	A=900
A =950		$A := A - temp$	B=2000
B=2000	write (A)	write (A)	
B=2000+50=2050	read (B)	read (B)	
B=2050	$B := B + 50$		
	write (B)		2000 +100=2100
	commit		B=2100
		$B := B + temp$	
		write (B)	
		commit	

How to check  
which kind of  
concurrent  
Schedule  
maintain  
Consistency

Final  $A=950$  and  $B=2100$ , now  $A+B =3050$  which violates Consistency

# Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus **serial execution** of a set of transactions **preserves database consistency**.
- A (possibly **concurrent**) **schedule is serializable** if it is **equivalent to a serial schedule**.
- Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**

# ***Simplified view of transactions***

- We **ignore operations** other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our simplified schedules consist of only **read** and **write** instructions.

# Conflicting Instructions

- Assume a schedule  $S$  in which there are two consecutive instructions,  $I$  and  $J$ , of transactions  $T_i$  and  $T_j$ , respectively ( $i \neq j$ ).
  - ▶ If  $I$  and  $J$  refer to different data items  
then we can swap  $I$  and  $J$  without affecting the results of any instruction
- In the schedule. However,
  - ▶ if  $I$  and  $J$  refer to the same data item  $Q$ , then the order of the two steps may matter.
- 1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict. (can be swapped)
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict. (can't be swapped)
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict (can't be swapped)
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict (can't be swapped)
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

# Conflict Serializability

- If a schedule **S** can be transformed into a schedule **S'** by a series of swaps of non-conflicting instructions, we say that **S** and **S'** are **conflict equivalent**.
- We say that a schedule **S** is **conflict serializable** if it is **conflict equivalent** to a serial schedule

$T_1$	$T_2$	$T_1$	$T_2$
read (A) write (A)	read (A) write (A)	read (A) write (A) read (B) write (B)	read (A) write (A)
read (B) write (B)	read (B) write (B)		read (B) write (B)
<b>S</b>		<b>S'</b>	

**S** is conflict serializable

# Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions. Therefore **Schedule 3 is conflict serializable**.

$T_1$	$T_2$
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6

# Conflict Serializability (Cont.)

- Example of a schedule that is **not conflict serializable**: **(schedule 7)**

$T_3$	$T_4$
read ( $Q$ )	
	write ( $Q$ )
write ( $Q$ )	

- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .



# Example

T1	T2
Read(A)	
	Read(A)
Read(B)	
Write(B)	
	Read(D)
	Write(D)
Write (A)	
	Read(C)
	Write(A)

**Is it Conflict Serializable?**

# Concurrency Control

# Concurrency Control

- ❑ Lock-Based Protocols
- ❑ Timestamp-Based Protocols
- ❑ Validation-Based Protocols
- ❑ Multiple Granularity
- ❑ Multiversion Schemes
- ❑ Insert and Delete Operations
- ❑ Concurrency in Index Structures

One of the fundamental properties of a **transaction is isolation**. When several transactions execute concurrently in the database, however, the isolation property **may no longer be preserved**.

To ensure that it is, the system must control the interaction among the concurrent transactions; this control is achieved through one of a variety of mechanisms called *concurrency control* schemes.

# Lock-Based Protocols

- One way to **ensure serializability** is to require that data items be **accessed in a mutually exclusive manner**
- A lock is a mechanism to **control concurrent access** to a data item
- Data items can be locked in two modes :
  - 1. **exclusive (X) mode**. If a transaction  $T_i$  has obtained an **exclusive-mode lock** (denoted by X) on item Q, then  $T_i$  can both **read** and **write Q**.
  - 2. **shared (S) mode**. If a transaction  $T_i$  has obtained a **shared-mode lock** (denoted by S) on item Q, then  $T_i$  can **read**, but **cannot write, Q**.
- Lock requests are **made to concurrency-control manager**.
- Transaction can proceed only after request is granted.

# Lock-Based Protocols (Cont.)

## □ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
  - but if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.

Let  $\{T_0, T_1, \dots, T_n\}$  be a set of transactions participating in a schedule  $S$ . We say that  $T_i$  **precedes**  $T_j$  in  $S$ , written  $T_i \rightarrow T_j$ , if there exists a data item  $Q$  such that  $T_i$  has held lock mode **A (S or X)** on  $Q$ , and  $T_j$  has held lock mode **B (S or X)** on  $Q$  later, and  $\text{comp}(A, B) = \text{false}$ .

If  $T_i \rightarrow T_j$ , then that precedence **implies** that in **any equivalent serial schedule**,  $T_i$  must appear before  $T_j$ .

# Lock-Based Protocols (Cont.)

- Example of a transaction performing locking (serially  $(T_1, T_2)$  OR  $(T_2, T_1)$ ) :

$T_1$ : lock-X( $B$ );  
read( $B$ );  
 $B := B - 50$ ;  
write( $B$ );  
unlock( $B$ );  
lock-X( $A$ );  
read( $A$ );  
 $A := A + 50$ ;  
write( $A$ );  
unlock( $A$ ).

$T_2$ : lock-S( $A$ );

read ( $A$ );

unlock( $A$ );

lock-S( $B$ );

read ( $B$ );

unlock( $B$ );

display( $A+B$ )

Assume  $A=100$  &  $B=200$

- If the transactions are executed serially, either in the order  $T_1, T_2$  or the
- order  $T_2, T_1$ , then transaction  $T_2$  will display the **value \$300**. No Loss of Consistency.
- If, however, these transactions are executed concurrently, (see schedule 1 next slide) then in this case, transaction  $T_2$  displays \$250, which is incorrect. The reason for this mistake is that the transaction  $T_1$  unlocked data item  $B$  too early, as a result of which  $T_2$  saw an inconsistent state.
- if  $A$  and  $B$  get updated in-between the read of  $A$  and  $B$ , the displayed sum would be wrong.

Assume A=100 & B=200

# Schedule 1-Concurrent

If, however, these transactions are executed concurrently, then schedule 1

$T_1$	$T_2$	concurrency-control manager
lock-x (B)  read (B) $B := B - 50$ write (B) unlock (B)	     lock-s (A)  read (A) unlock (A) lock-s (B)  read (B) unlock (B) display (A + B)	grant-x (B, $T_1$ )     grant-s (A, $T_2$ )  grant-s (B, $T_2$ )   grant-x (A, $T_2$ )
lock-x (A)  read (A) $A := A + 50$ write (A) unlock (A)		

In this case,  
transaction  $T_2$   
displays \$250,  
which is incorrect.

Reason is - $T_1$  unlocked  
data item B too early.

What if hold lock for  
longer duration ?

A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.

# Pitfalls of Lock-Based Protocols (Deadlock)

- Consider the partial schedule

$T_3$	$T_4$
lock-x ( $B$ )	
read ( $B$ )	
$B := B - 50$	
write ( $B$ )	
	lock-s ( $A$ )
	read ( $A$ )
	lock-s ( $B$ )
lock-x ( $A$ )	

locking can lead to an undesirable situation- **Deadlock**

- Neither  $T_3$  nor  $T_4$  can make progress — executing **lock-S( $B$ )** causes  $T_4$  to wait for  $T_3$  to release its lock on  $B$ , while executing **lock-X( $A$ )** causes  $T_3$  to wait for  $T_4$  to release its lock on  $A$ .
- Such a situation is called a **deadlock**.
  - To handle a deadlock one of  **$T_3$  or  $T_4$  must be rolled back** and its locks released.



# Pitfalls of Lock-Based Protocols (Deadlock)

If we do not use locking, or if we **unlock data items as soon as possible** after reading or writing them, we may get **inconsistent states**.

On the other hand, if we **do not unlock a data item before requesting a lock on another data item**, **deadlocks may occur**.

**Which one to Choose ? Deadlock or getting into inconsistent state**

**Deadlocks** are definitely **preferable to inconsistent states**, since they can be handled by rolling back of transactions, whereas

**Inconsistent states** may **lead to real-world problems** that cannot be handled by the database system.

We shall require that each transaction in the system follow a set of rules, called a **locking protocol**, indicating when a transaction may lock and unlock each of the data items.

Locking protocols **restrict the number of possible schedules**. The set of all such schedules is a **proper subset of all possible serializable schedules**.

# Pitfalls of Lock-Based Protocols (Cont.)

- The potential for **deadlock** exists in most locking protocols. **Deadlocks are a necessary evil.**
- **Starvation** is also possible if concurrency control manager is badly designed. For example:
  - A transaction may be **waiting ( $T_1$ ) for an X-lock on an item**, while a sequence of **other transactions ( $T_2, T_3, \dots$ ) request and are granted an S-lock** on the same item.
  - The same transaction ( $T_1$ ) is **repeatedly rolled back due to deadlocks**.
- Concurrency control manager can be designed to prevent starvation
- **We can avoid starvation of transactions by granting locks in the following manner:**
- When a transaction  $T_i$  **requests a lock on** a data item  $Q$  in a particular **mode  $M$** , the concurrency-control manager grants the lock provided that:
  - **1.** There is **no other** transaction **holding a lock on  $Q$  in a mode that conflicts with  $M$** .
  - **2.** There is **no other** transaction that is **waiting** for a lock on  $Q$  and that made its lock request before  $T_i$ .

# The Two-Phase Locking Protocol

- This is a protocol which ensures conflict-serializable schedules.
- **Phase 1: Growing Phase**
  - transaction **may obtain** locks
  - transaction **may not release** locks
- **Phase 2: Shrinking Phase**
  - transaction **may release** locks
  - transaction **may not obtain** locks
- The protocol assures serializability.
- The two-phase locking protocol ensures conflict serializability.
- Consider any transaction. The point in the schedule where the transaction has obtained its **final lock** (the **end of its growing phase**) is called the **lock point** of the transaction. Now, transactions can be **ordered according to their lock points** this ordering is, in fact, a serializability ordering for the transactions.

# Pitfall in 2-phase Locking

In addition to being serializable, schedules should be cascade less. **Cascading rollback may occur under two-phase locking.**

consider the partial schedule given below-

$T_5$	$T_6$	$T_7$
lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A)	lock-X(A) read(A) write(A) unlock(A)	lock-S(A) read(A)

failure of  $T_5$   
after the  
read(A)  
step of  $T_7$

Each transaction observes the two-phase locking protocol, but the **failure** of  $T_5$  **after the read(A) step of  $T_7$**  leads to **cascading rollback** of  $T_6$  and  $T_7$ .

# The Two-Phase Locking Protocol (Cont.)

- ❑ Two-phase locking *does not ensure freedom from deadlocks*
  - ❑ ( see the schedule containing T3 & T4 slide 7)
- ❑ Cascading roll-back is possible under two-phase locking.
- ❑ To avoid this, follow a modified protocol called strict two-phase locking.
  - ❑ Here a transaction must **hold all its exclusive locks** till it commits/aborts.
- ❑ This requirement ensures that any data written by an uncommitted transaction are **locked in exclusive mode until the transaction commits**, preventing any other transaction from reading the data.
- ❑ **Rigorous two-phase locking** is even stricter: here all locks are **held till commit/abort**. In this protocol transactions can be serialized in the order in which they commit.

**END**